

Curiosity-driven exploration for PPO agents in sparse reward environments




[1]

Motivation

- Learning an agent for low-dimensional state spaces has become a trivial task
- Many environments are openly available with OpenAIs Gym
- Proximal Policy Optimization is one of the most popular methods [\[5\]](#)
 - due to stability and performance and versatility
 - Still has limitations in sparse reward environments and deep exploration tasks
- To enhance the agent to explore and discover novel states, intrinsic rewards could offer a solution
 - New or distant rewards could be discovered which can lead to faster learning
 - Faster learning can lead to better overall performance with the same interaction budget
- Random Network Distillation offers the potential enhancement for exploration behavior
- Crafter is a suitable environment to test this method, because of its complexity and deep exploration challenges

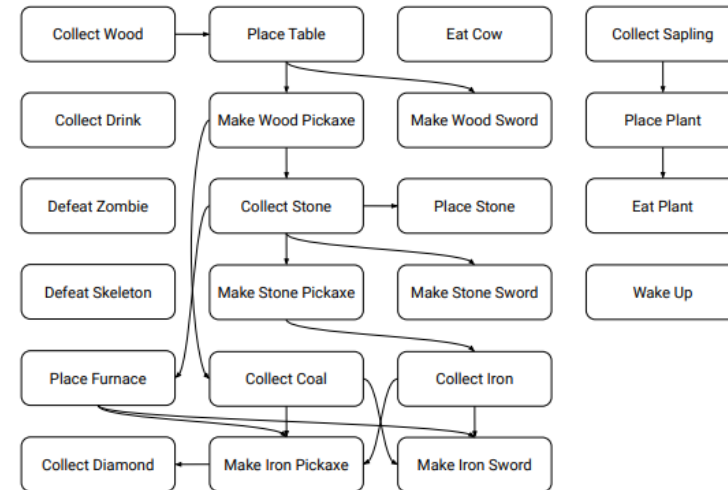
Related Work

- Crafter with roguelike elements: in Jax Craftax [\[2\]](#)
- VizDoom: Open-Source adaptation of the original Doom for AI research [\[3\]](#)
 - Which was could be solved by enhancing exploration
- PPO - RND for low-complexity state space games [\[6\]](#)



Crafter Environment

- 2D abstraction of Minecraft
- 17 possible actions (e.g., move (W, A, S, D))
- 22 Achievements ~ 22 max score
- Map gets generated (64,64,3)
- 4 levels of survival (HP, food, water, etc.)
- Needs to collect resources for crafting
- Creatures can attack at night
- -0.1 reward for damage, most episodes end on .1 because the agent died
- Episodes ends after 10.000 steps



[1]



PPO

- Improves the decision-making by optimizing an agents policy through interactions [\[5\]](#)
- Adjusts the agents policy based on the rewards that are accumulated during episodes
- Clipping - penalty for policy updates that diverge highly compared to the old policy
 - The probability ratio $r_t(\theta)$ between the updates is limited with $1 - \epsilon, 1 + \epsilon$
 - Idea comes from Kullback-Leibler divergence which is used in PPOs predecessor TRPO [\[7\]](#).
- Entropy term balances exploration behavior
- Advantage function A_t provides the probability of which action should be taken by subtracting the Q-value with the output of the value-function

$$L^{\text{CLIP}}(\theta) = \mathbb{E}_t[\min(r_t(\theta)A_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon)A_t)]$$

RND

- Starts from randomly initialized untrained neural network
- Over time distill it into a trained network - Network is reduced to a smaller size
- Initially trained on Montezumas Revenge which requires the agent to work towards late rewards with collected keys
 - Which aligns with Crafters structure
- Generates intrinsic rewards that are based on the prediction error
 - Uses fixed target network and a predictor network to calculate error
 - High prediction error: Agent does not know the state → high intrinsic reward for state novelty
- Learning process: intrinsic rewards + reward structure from the environment
- Useful for sparse reward environments and efficient exploration

Approach: PPO with RND

- PPO uses Generalized Advantage Estimation (GAE)
 - Controls bias-variance trade-off and uses TD-error
 - In the project the GAE is calculated once for external and intrinsic rewards
 - Both results are summed and weighted with fix coefficients in favour of the extrinsic reward

$$A_t^{\text{GAE}} = \sum_{l=0}^{\infty} (\gamma\lambda)^l \delta_{t+l}$$

$$A_{\text{internal}} = \sum_{t=0}^{\infty} \left(r_t^{\text{int}} + \gamma V_{\text{int}}(s_{t+1}) - V_{\text{int}}(s_t) \right)$$

$$A_{\text{external}} = \sum_{t=0}^{\infty} \left(r_t^{\text{ext}} + \gamma V(s_{t+1}) - V(s_t) \right)$$


$$A_t = \alpha * A_{\text{external}} + \beta * A_{\text{internal}}$$

Approach: Algorithm flow

1. Collect experience from environment
2. Compute extrinsic rewards (milestones)
3. Compute intrinsic rewards (prediction error with unvisited states)
 - RND generates intrinsic rewards based on unvisited states
 - Aimed to balance task completion and exploration
4. Combine both rewards and update PPO policy
 - With modified Advantage calculation
5. Update RND network periodically for better exploration behavior
 - Based on novelty

Experiments and Results

Experiment Setup: 500 episodes, 5 seeds



PPO: After approx. 50 episodes the agent reaches the total reward of 4 and has a robust overall training performance

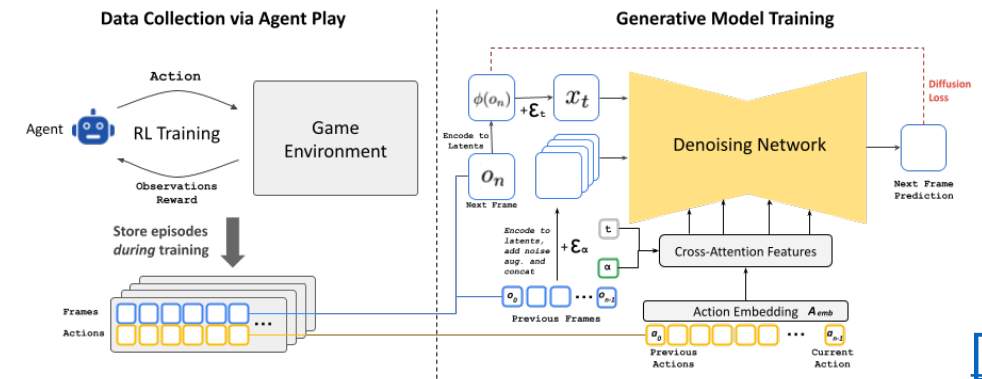
PPO + RND: Agent reaches score of 4 faster, but can not keep the performance, hence it is a less robust performance and has a higher variance in achieved milestones.

Discussion

- PPO has overall the better performance, due to robustness and mean score
- PPO + RND shows that intrinsic rewards can lead to faster rewards
 - But instead of exploiting that, the agent explores in the following episodes
 - The configuration of the hyperparameters has a major impact on the result, but some lack in stability
- **Possible improvements:**
 - Find a better configuration for hyperparameters
 - Re-balance the weights of each advantage
 - Use RND with a time limit (set coefficient to zero after n-steps)

Future Work

- Inspect each milestone separately to see how each method impacts
- It can be tested on similar games like Craftax [\[2\]](#), which has less sparse rewards and a rogue-like game structure
 - Which makes exploration more impactful
- If enough compute is available, test both methods on Minecraft
- Test each method in GameNGen [\[8\]](#) to see impact of exploration on real-time generated game content to measure playability
 - Is PPO with RND a better benchmark for human replay then PPO?



Questions?

References

- [1] Hafner, D. (2021). Benchmarking the Spectrum of Agent Capabilities. arXiv preprint:2109.06780
- [2] Matthews, M. et al. (2024). Craftax: A Lightning-Fast Benchmark for Open-End Reinforcement Learning. International Conference on Machine Learning (ICML).
- [3] Marek Wydmuch, Michał Kempka, and Wojciech Jas´kowski. ViZDoom Competitions: Playing Doom from Pixels. *IEEE Transactions on Games*, 11(3):248–259, 2019. doi: 10.1109/TG.2018.2877047. The 2022 IEEE Transactions on Games Outstanding Paper Award.
- [4] Bradbury, J. et al (2018). JAX: composable transformations of {P}ython+{N}um{P}y programs. <http://github.com/google/jax>.
- [5] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. arXiv preprint arXiv:1707.06347, 2017.
- [6] Nugroho, W. (2021). Reinforcement Learning PPO RND [Source code]. GitHub. https://github.com/wisnunugroho21/reinforcement_learning_ppo_rnd
- [7] Schulman et al. (2015). Trust Region Policy Optimization. arXiv:1502:05477.
- [8] Dani Valevski, Yaniv Leviathan, Moab Arar, and Shlomi Fruchter. Diffusion models are real-time game engines, 2024. URL <https://arxiv.org/abs/2408.14837>.